



TensorFlow tutorial

Adam Elwood

1st Terascale School of Machine Learning, Oct 2018

Introduction and outline

- This is a tutorial designed to give a very basic introduction to the main concepts behind TensorFlow, consisting of:
 - A very brief overview of neural networks from a TensorFlow standpoint
 - TensorFlow as a computational graph framework
 - How Automatic Differentiation works
 - A simple example of classifying two different Gaussians with a neural network
- Following this there is an exercise to practice writing your own simple network with TensorFlow in a HEP context
- Thanks to Stefan Wunsch who provided the inspiration and lots of material for this tutorial in the IML workshop in April
 - <https://indico.cern.ch/event/668017/contributions/2947042/>
 - https://github.com/stwunsch/iml_tensorflow_keras_workshop

Introduction to TensorFlow

- TensorFlow is a low-level implementation of operations needed to implement neural networks in multi-threaded CPU and multi GPU environments
- Differences with **PyTorch** (Facebook vs Google)
 - TensorFlow designed to use static graph definition, you have to build the graph then compile it before running
 - PyTorch allows the graph to be built and executed dynamically
 - TensorFlow is currently the most widely used but PyTorch is gaining popularity
 - Generally speaking: TensorFlow better for production code and PyTorch better for research
 - More info: <https://towardsdatascience.com/pytorch-vs-tensorflow-spotting-the-difference-25c75777377b>
 - **Disclaimer:** these frameworks are developing so quickly this information may not stay relevant
- **Keras** provides high-level convenience wrapper for backend libraries, predominantly TensorFlow, to implement neural network models
 - Will be covered in much more detail later

PYTORCH

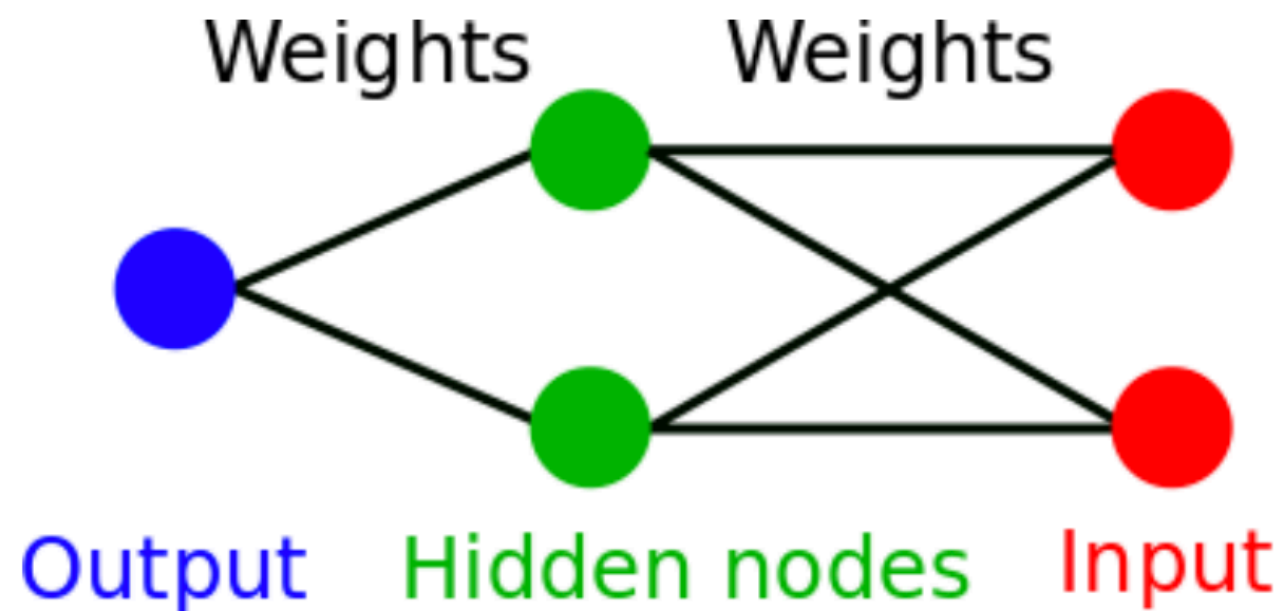
vs



Source for code

- Github repository containing the slides and examples:
 - https://github.com/aelwood/iml_tensorflow_keras_workshop
- Notebooks relevant to this tutorial in 'tensorflow' folder
- Setup instructions:

A simple neural network



Neural Network: $f(x)$

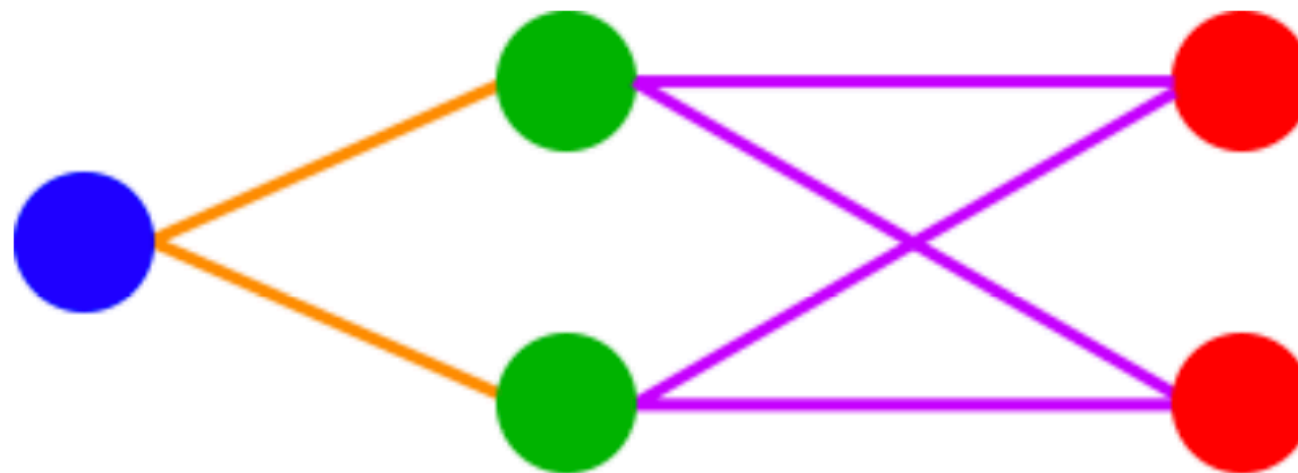
- ▶ **Important:** A neural network is only a mathematical function. No magic involved!
- ▶ **Training:** Finding the best function for a given task, e.g., separation of signal and background.

Mathematical representation

- **Why do we need to know this?**

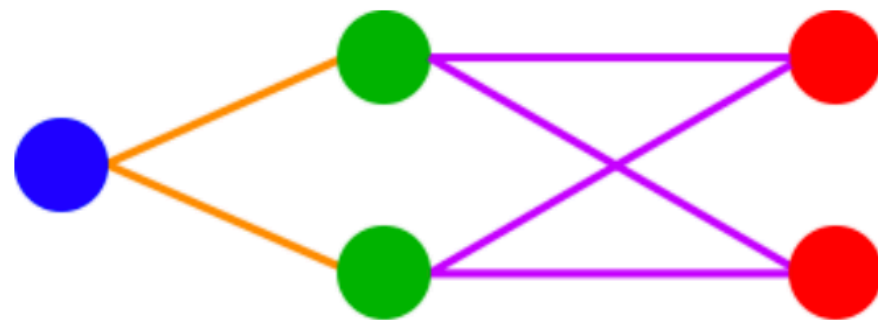
- TensorFlow implements these mathematical operations explicitly.

- Basic knowledge to understand Keras' high-level layers.



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

Mathematical representation (2)



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

$$\text{Input : } x = \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix}$$

$$\text{Weight : } W_1 = \begin{bmatrix} W_{1,1} & W_{1,2} \\ W_{2,1} & W_{2,2} \end{bmatrix}$$

$$\text{Bias : } b_1 = \begin{bmatrix} b_{1,1} \\ b_{2,1} \end{bmatrix}$$

Activation : $\sigma(x) = \tanh(x)$ (as example)

Activation is applied elementwise!

The “simple” neural network written as full equation:

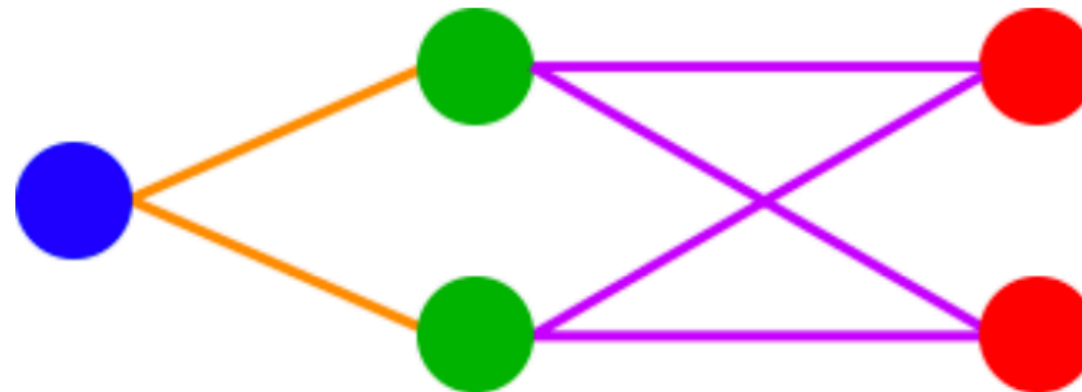
$$f_{\text{NN}} = \sigma_2 \left(\begin{bmatrix} b_{1,1}^2 \end{bmatrix} + \begin{bmatrix} W_{1,1}^2 & W_{1,2}^2 \end{bmatrix} \sigma_1 \left(\begin{bmatrix} b_{1,1}^1 \\ b_{2,1}^1 \end{bmatrix} + \begin{bmatrix} W_{1,1}^1 & W_{1,2}^1 \\ W_{2,1}^1 & W_{2,2}^1 \end{bmatrix} \begin{bmatrix} x_{1,1} \\ x_{2,1} \end{bmatrix} \right) \right)$$

What is TensorFlow?

TensorFlow is an open source software library for **numerical computation using data flow graphs.**

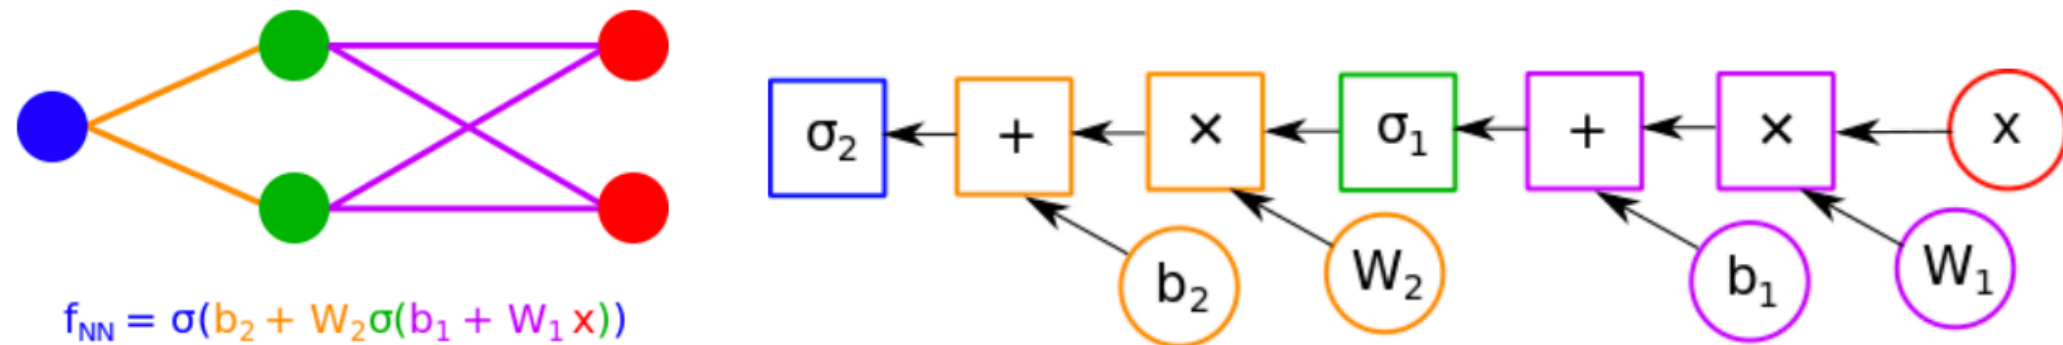
Nodes in the graph represent mathematical operations, while the graph edges represent the multidimensional data arrays (tensors) communicated between them.

- ▶ **In first place:** TensorFlow is not about neural networks.
- ▶ But it is a **perfect match** to implement neural networks efficiently!



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$

Computational graphs



Example neural network



According computational graph

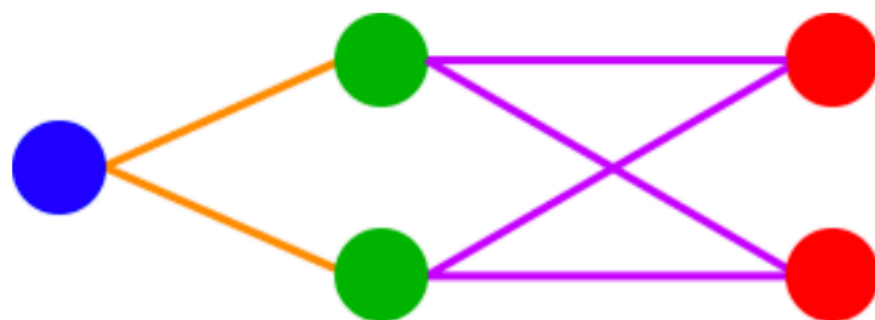
- ▶ TensorFlow implements all needed **mathematical operations for multi-threaded CPU and multi GPU** environments.
- ▶ Computation of neural networks using data flow graphs is a perfect match!

***TensorFlow** is an open source software library for numerical computation using data flow graphs. **Nodes** in the graph **represent mathematical operations**, while the **graph edges represent the multidimensional data arrays (tensors)** communicated between them.*

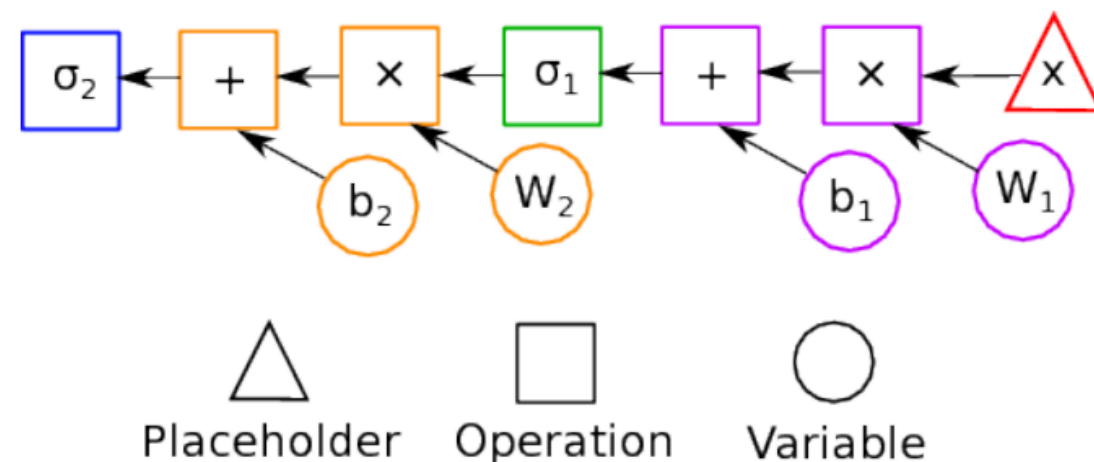
Basic blocks to build graphs in TensorFlow

► Basic blocks:

- **Placeholders:** Used for injecting data into the graph, e.g., the inputs x of the neural network
- **Variables:** Free parameters of the graph, e.g., the weight matrices W of the neural network
- **Operations:** Functions that operate on data in the graph, e.g., the matrix multiplication of W_1 and x



$$f_{\text{NN}} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$



Run the graph in a TensorFlow session

- ▶ A **graph** in TensorFlow can be run inside a **session**.
- ▶ Following example calculates $y = W \cdot x$ using TensorFlow:

Computational graph:

$$y = W \cdot x = \begin{pmatrix} 1 & 2 \end{pmatrix} \cdot \begin{pmatrix} 3 \\ 4 \end{pmatrix} = 11$$

TensorFlow code:

```
import tensorflow as tf
import numpy as np
```

```
# Build the graph y = W * x
```

```
x = tf.placeholder(tf.float32) # A placeholder
```

```
W = tf.get_variable("W", initializer=[[1.0, 2.0]]) # A variable
```

```
y = tf.matmul(W, x) # An operation
```

```
with tf.Session() as sess: # The session
```

```
    sess.run(tf.global_variables_initializer()) # Initialize variables
```


```
    result = sess.run(y, feed_dict={x: [[3.0], [4.0]]}) # Run graph
```

See example:
tensorflow/
xor.ipynb

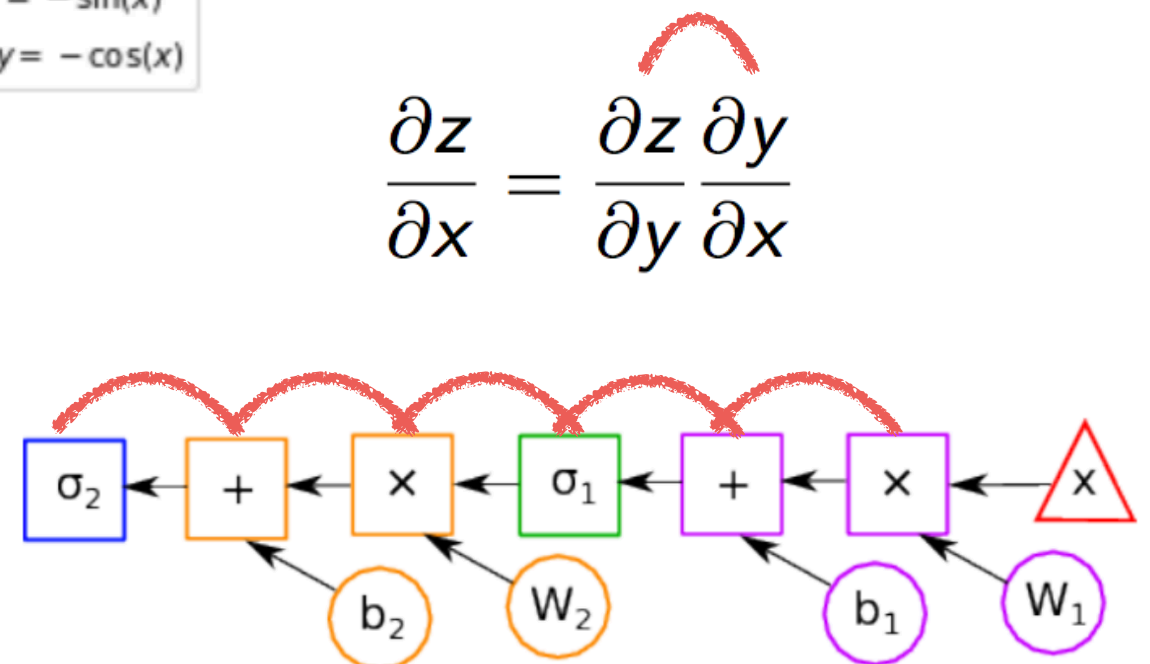
Automatic differentiation

- During training of neural networks make an excessive use of gradients during optimisation
- (Almost) each operation in TensorFlow is shipped with an inbuilt gradient
- Computation of full gradient using the chain-rule of derivatives through the computational graph
- Explicit TensorFlow call: `tensorflow.gradients(z, x)`

See example: [tensorflow/automatic_differentiation.ipynb](#)

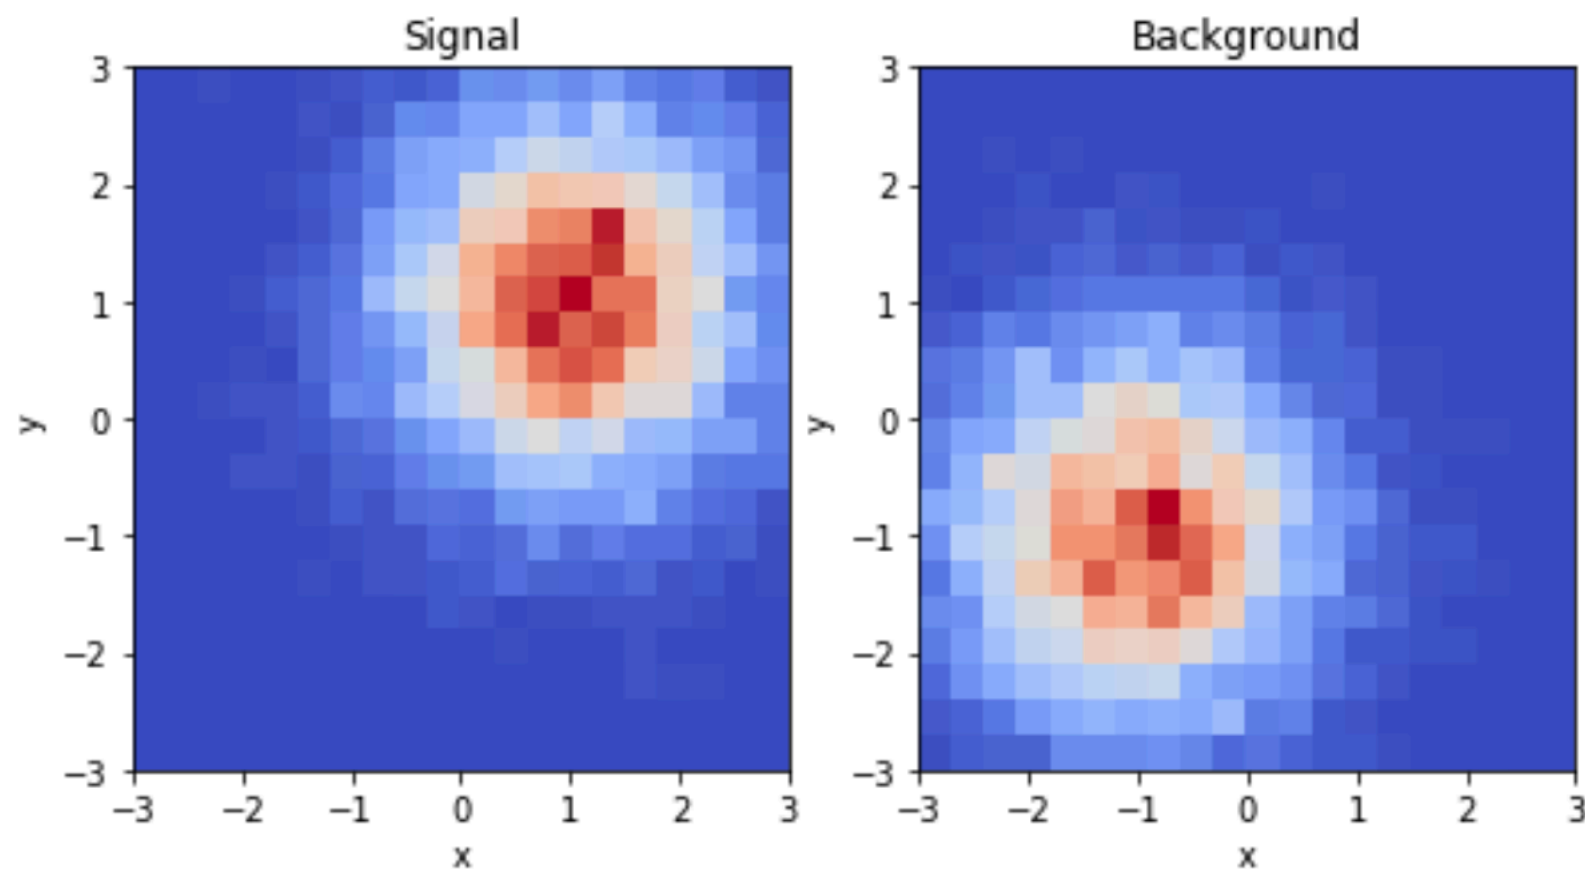


$$\frac{\partial z}{\partial x} = \frac{\partial z}{\partial y} \frac{\partial y}{\partial x}$$



Concrete TensorFlow example: classification of toy Gaussian models

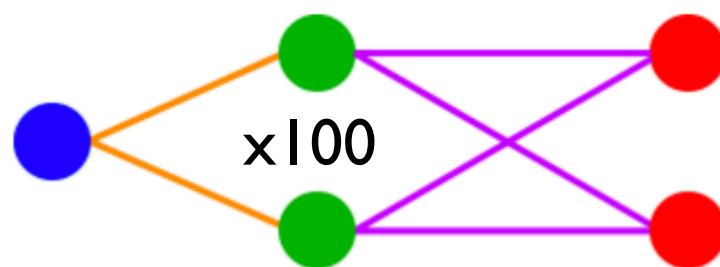
- Carry out a basic classification task separating signal from background
- We will build a single hidden layer neural network with tensor flow and optimise with the Adam gradient descent algorithm




See example: [tensorflow/gaussian.ipynb](#)

Define the network architecture

```
In [4]: def model(x):  
    with tf.variable_scope("model") as scope:  
        w1 = tf.get_variable('w1', shape=(2, 100), dtype=tf.float64,  
                             initializer=tf.random_normal_initializer(stddev=0.1))  
        b1 = tf.get_variable('b1', shape=(100), dtype=tf.float64,  
                             initializer=tf.constant_initializer(0.1))  
        w2 = tf.get_variable('w2', shape=(100, 1), dtype=tf.float64,  
                             initializer=tf.random_normal_initializer(stddev=0.1))  
        b2 = tf.get_variable('b2', shape=(1), dtype=tf.float64,  
                             initializer=tf.constant_initializer(0.1))  
  
        l1 = tf.nn.relu(tf.add(b1, tf.matmul(x, w1)))  
        logits = tf.add(b2, tf.matmul(l1, w2))  
    return logits, tf.sigmoid(logits)  
  
x = tf.placeholder(tf.float64, shape=[None, 2])  
logits, f = model(x)
```




$$f_{NN} = \sigma(b_2 + W_2 \sigma(b_1 + W_1 x))$$



Define the training

- Train using the binary cross entropy calculated on logits
 - Cross entropy usually best choice for classification
 - Equivalent to minimising the negative log likelihood
- Output with a sigmoid activation function
 - Sigmoid trained to 1 for signal and 0 for background
- Use the Adam optimiser: gradient descent with momentum and path correction

$$C = -\frac{1}{n} \sum_x [y \ln a + (1 - y) \ln(1 - a)],$$



```
In [5]: labels = tf.placeholder(tf.float64, shape=[None, 1])
loss = tf.reduce_mean(
    tf.nn.sigmoid_cross_entropy_with_logits(labels=labels, logits=logits))
minimize_loss = tf.train.AdamOptimizer().minimize(loss)
```

Compile and run the model

- Everything is done within a TensorFlow session
 - Define this session and call a global variable initialisation
- Pass the predefined loss and loss minimisation procedure to the session, broadcasting the input variables
- Additionally calculate the loss on the test data

```
In [6]: sess = tf.Session()
sess.run(tf.global_variables_initializer())

loss_train = []
loss_val = []
for i_step in range(100):
    loss_, _ = sess.run([loss, minimize_loss],
                        feed_dict={x: data_train, labels: labels_train})
    loss_train.append(loss_)

    loss_ = sess.run(loss, feed_dict={x: data_val, labels: labels_val})
    loss_val.append(loss_)
```


Extra TensorFlow features

- TensorFlow is designed to perform highly-efficient computations and ships many useful features
 - **Data-loading** often bottleneck if not all data fits in memory (very common for image processing!)
 - TensorFlow provides input pipelines directly inbuilt in the graph
 - Full utilisation of CPU/GPU by loading data from disk in queues in memory concurrently
- **'TensorBoard'** can be used for visualisation of graphs
- **'Eager execution'** recently added to emulate PyTorch style
 - <https://www.tensorflow.org/guide/eager>
- **Keras** provides an excellent high layer wrapper, making development very convenient
- Many more features not mentioned here...

See example: tensorflow/
queues.ipynb

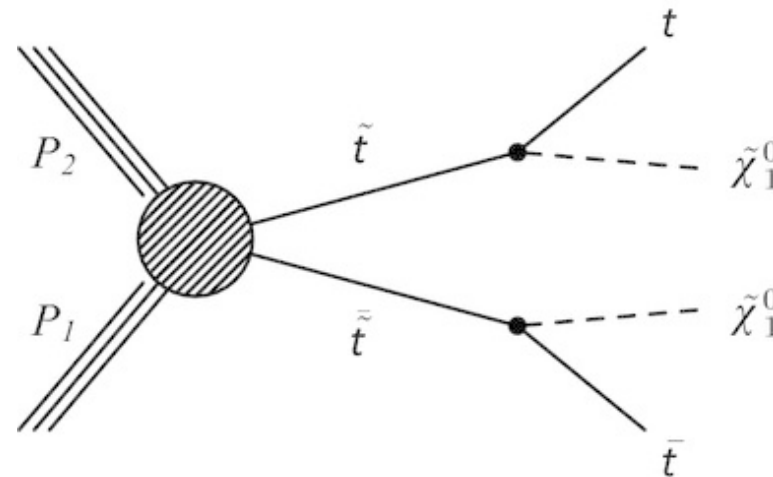
Further reading

- Stanford course on TensorFlow
 - Very well done and highly entertaining course!
 - Lecturer working in the field (OpenAI, DeepMind, Google, ...)
 - Small Keras part held by Francois Chollet (author of Keras!)
 - Link: <https://web.stanford.edu/class/cs20si/syllabus.html>
- Free textbook written by Ian Goodfellow, Yoshua Bengio and Aaron Courville:
 - Leading figures in current machine learning research
 - Covers much of what you could want to know
 - Link: <http://www.deeplearningbook.org/>

Practice with a physics example

- Now you can try to write your own TensorFlow code in a toy physics scenario
- A background vs signal classification task for a stop SUSY model (with a top anti-top standard model background)

$$m_{\text{stop}} = 600 \text{ GeV},$$
$$m_{\text{LSP}} = 400 \text{ GeV}$$



- Try to build a two layer network and train with mini-batch optimisation
- Instructions and skeleton code in `tensorflow/physicsExample.ipynb`
 - A solution available if you are stuck `tensorflow/physicsExample_solution.ipynb`
- After this task, have a go at doing a regression task, also in the notebook